1. Current setup: manual installers

MUNGBEANS / RunPod reference / Updated 23 October 2026 (Perth)

Current plan: startup fetches installers only, then runs the image normal /start.sh so Jupyter can open. Models and custom nodes install only when you run a numbered command. This replaces the earlier all-in-one startup instructions.

Container start command: one field

```
/bin/bash -c '
mkdir -p /workspace
f=/workspace/fetch-installers.sh
if curl -fLsS --connect-timeout 15 --max-time 60 --retry 2 \
  https://raw.githubusercontent.com/martynguyen420-debug/MUNGBEAN/main/runpod-startup/fetch-installers.sh \
  -o "$f.new" && bash -n "$f.new"; then
  mv "$f.new" "$f"
fi
if [ -s "$f" ]; then
  timeout 180 bash "$f" || echo "Installer fetch failed; continuing normal startup"
fi
exec /bin/bash /start.sh
'
```

For one-line paste, use **manual-template-command-single-field.txt** in GitHub. If the editor has separate entrypoint and command fields, use **manual-template-start.json**. Requires the existing image to provide /start.sh; first pod startup remains unverified.

In Jupyter: Terminal

```
cd /workspace
ls 1 2 3 4 5 6 7
bash 6
```

6 custom nodes first. Then **bash 1** Phr00t v19 + v23; **bash 2** Klein; **bash 3** Miraclein; **bash 4** Aisha; **bash 5** KreaNoct. Then **bash 7** starts missing 8188/8189 services and preserves existing services. Expose both ports as HTTP. Include bash before each number.

Existing installation and MUNG

Default root: /workspace/runpod-slim/ComfyUI. Python: its .venv-cu128/bin/python or the sibling environment. Confirm these on the pod. No automatic core/CUDA/torch upgrade. Set HF_TOKEN privately and accept the Klein license if required.

MUNG: each remote session asks for the current pod 8188 URL; the old saved pod address is not reconnected automatically. The new URL is saved only after validation. Fixed desktop source is in GitHub mung-desktop and the F: recovery folder.

2. Models available by number

MUNGBEANS / RunPod reference / Updated 23 October 2026 (Perth)

12 locked files / 134.05 GB (124.85 GiB), excluding ComfyUI, existing models and outputs. Allow at least 150 GB free for the model set, plus space for generations.

Model	Required support / loader
Phr00t Rapid v19 + v23 Q8_0 GGUF	Qwen2.5-VL 7B Q8_0 GGUF + matching BF16 mmproj; Qwen Image VAE. UnetLoaderGGUF and CLIPLoaderGGUF type qwen_image.
FLUX.2 Klein 9B distilled BF16	Qwen3 8B BF16 encoder + Flux2 VAE. Core UNETLoader; CLIPLoader type flux2. Exact-byte ModelsLab mirror; original BFL license applies.
Miraclein NSFW v3.0 FP8 9B	Shares Klein Qwen3 8B encoder and Flux2 VAE. Core UNETLoader. Exact named file from EllaPriest45 mirror.
Aisha NSFW beta v9.7 distilled BF16	Shares Klein Qwen3 8B encoder and Flux2 VAE. Downloaded from the official Aisha dataset repository.
KreaNoct V4 INT8 ConvRot	Qwen3-VL 4B FP8-scaled encoder + shared Qwen Image VAE. Core diffusion-model loader. Publisher requires ComfyUI >=0.37.

Important workflow differences

Phr00t GGUF is loaded as a diffusion model with separate encoder and VAE. Do not use the AIO safetensors checkpoint-loading commands from the old October 13 guide for this GGUF setup. The mmproj stays beside its encoder with the exact name in the manifest.

KreaNoct: publisher recommends 8 steps, CFG 1, euler/simple. Phr00t v23: euler_ancestral/beta, CFG 1, 4-8 steps. Miraclein mirror filename suggests 12 steps/euler/CFG 1.1. Use a matching workflow for Aisha and the official distilled Klein workflow for Klein.

Custom nodes and compatibility

Installs only pinned ComfyUI-GGUF, matching the minimal recovered node installer. Existing unrelated nodes are preserved. No automatic core, torch, CUDA or existing package upgrade. Only missing small GGUF dependencies are installed without dependency resolution.

Setup checks only the selected model family: Phr00t/Qwen, Klein or Krea2. KreaNoct requires ComfyUI 0.37 or newer; the template screenshot showed 0.30.0. Missing support is reported; the core is not silently updated. Model pairing was checked from repository records, the GGUF loader source and relevant tensor headers. Actual target-GPU inference still needs testing.

3. Check startup and recover

MUNGBEANS / RunPod reference / Updated 23 October 2026 (Perth)

Logs and serving checks

```
cat /workspace/.runpod-startup/group-6-status.json
cat /workspace/.runpod-startup/group-1-locations.json
curl --connect-timeout 5 --max-time 15 -fsS \
  http://127.0.0.1:8188/system_stats
curl --connect-timeout 5 --max-time 15 -fsS \
  http://127.0.0.1:8188/object_info/UnetLoaderGGUF
curl --connect-timeout 5 --max-time 15 -fsS \
  http://127.0.0.1:8188/object_info/CLIPLoaderGGUF
```

Startup fetches scripts only. Run numbered installers in Jupyter, then use your existing ComfyUI launcher for 8188. If nodes were installed while it was running, restart only that service. Leave 8189 alone. An installer failure is reported in the terminal and group status file; rerun after resolving the cause.

Persistent files and a different GPU

Mount the same persistent storage at the expected paths to reuse models on another compatible GPU pod. A new empty disk means a new download. Network volumes must be available in the chosen GPU data center.

Completed files are verified with size and SHA256, then tracked by a fingerprint receipt. Unchanged files are not downloaded or rehashed each launch. Run VERIFY_ALL=1 bash 1 (or another model number) to rehash that group.

Interrupted transfers resume from .part files. A mismatching completed file or corrupt full partial is preserved and reported. Inspect it against models.lock.json before moving it aside manually; automatic startup does not repeatedly replace it.

Instance paths and optional settings

```
COMFY_ROOT=/workspace/runpod-slim/ComfyUI
COMFY_PYTHON=/workspace/runpod-slim/ComfyUI/.venv-cu128/bin/python
STARTUP_REF=main
```

These optional environment settings select the existing installation and GitHub revision. Download locations are independent of the terminal current folder. paths.json rejects unapproved destinations; MODEL_LOCATIONS.md lists every full default path. No extra YAML is needed for ComfyUI native folders.

GitHub bundles are cached under /workspace/.mung-installers. An unavailable GitHub fetch falls back to the validated cached bundle. The first boot needs access to GitHub. STARTUP_REF can be a reviewed full Git commit SHA to pin the bundle.

Proof still required on the pod

Local checks passed Bash syntax and five real HTTP download fixtures: reuse, resume, bad-part preservation, corrupt-file rejection and HTTP 403 handling. No Linux/GPU workflow has been run for this bundle. After startup, open the 8188 proxy and generate a small test image with each model you intend to use.

4. Exact files and source record

MUNGBEANS / RunPod reference / Updated 23 October 2026 (Perth)

All files are locked to full repository commit revisions and SHA256 hashes in **models.lock.json**. The following is the download inventory; the machine-readable manifest is the authority for complete hashes.

diffusion_models/Qwen-Rapid-NSFW-v23_Q8_0.gguf

Novice25/Qwen-Image-Edit-Rapid-AIO-GGUF / 21.751 GB

diffusion_models/Qwen-Rapid-AIO-NSFW-v19_Q8_0.gguf

Novice25/Qwen-Image-Edit-Rapid-AIO-GGUF / 21.751 GB

text_encoders/Qwen2.5-VL-7B-Instruct-Q8_0.gguf

unsloth/Qwen2.5-VL-7B-Instruct-GGUF / 8.099 GB

text_encoders/Qwen2.5-VL-7B-Instruct-mmproj-BF16.gguf

QuantStack/Qwen-Image-Edit-GGUF / 1.354 GB

vae/qwen_image_vae.safetensors

Comfy-Org/Qwen-Image_ComfyUI / 0.254 GB

diffusion_models/flux-2-klein-9b.safetensors

ModelsLab/FLUX.2-klein-9B / 18.157 GB

text_encoders/qwen_3_8b.safetensors

Comfy-Org/flux2-klein-9B / 16.382 GB

vae/flux2-vae.safetensors

Comfy-Org/flux2-klein-9B / 0.336 GB

diffusion_models/aisha_nsfw_beta_v9_7_distilled_bf16.safetensors

Aisha-AI-Official/flux-2-klein-models / 18.157 GB

diffusion_models/Miraclein NSFW v3.0 FP8 - Klein9B - 12steps,euler,cfg1.1.safetensors

EllaPriest45/Klein9B_Checkpoints / 9.079 GB

diffusion_models/KreaNoct_V4_int8_convrot.safetensors

Noctaluna/Krea2-Noct-Uncensored / 13.493 GB

text_encoders/qwen3vl_4b_fp8_scaled.safetensors

Comfy-Org/Krea-2 / 5.242 GB

GitHub reference

github.com/martynguyen420-debug/MUNGBEAN / runpod-startup

Current entrypoint: fetch-installers.sh. Current config: manual-template-start.json and manual-template-command-single-field.txt. Installers: 1-6, install-group.sh, manual-install.py, setup.py. Paths and hashes: paths.json, MODEL_LOCATIONS.md, models.lock.json, nodes.lock.json. GitHub hosts only the scripts/configuration; model bytes come from the pinned Hugging Face repositories.

Aisha source: huggingface.co/datasets/Aisha-AI-Official/flux-2-klein-models. Miraclein source: huggingface.co/EllaPriest45/Klein9B_Checkpoints, a third-party Civitai backup. KreaNoct source: huggingface.co/Noctaluna/Krea2-Noct-Uncensored. Full sources and limitations are in README.md.

Updated from the newest RunPod commands PDF found on F: (13 October 2026). Original preserved. Recovered shell scripts provided the exact Aisha/Miraclein URLs. Publisher/API metadata and relevant headers rechecked on 22 October 2026.