1. Custom template startup

MUNGBEANS / RunPod reference / Updated 22 October 2026 (Perth)

Fetch the startup bundle from GitHub on every pod launch. Model files stay on persistent storage and are reused after verification. Starts ComfyUI on **8188**; leaves the separate **8189** service alone.

Entrypoint

```
/bin/bash -c
```

Command - paste the whole block

```
set -eu
mkdir -p /workspace
f=/workspace/runpod-template-bootstrap.sh
if curl -fLsS --connect-timeout 20 --max-time 90 --retry 3 \
  https://raw.githubusercontent.com/martynguyen420-debug/MUNGBEAN/main/runpod-startup/github-bootstrap.sh \
  -o "$f.new" && bash -n "$f.new"; then
  mv "$f.new" "$f"
else
  echo "GitHub unavailable: trying cached bootstrap"
fi
test -s "$f"
exec bash "$f"
```

If Command accepts one line only, use **template-command.txt** in the GitHub folder. For a JSON entrypoint/command editor, use **template-start.json**. Both contain the same startup behavior.

Required existing installation

```
/workspace/runpod-slim/ComfyUI/main.py
/workspace/runpod-slim/ComfyUI/.venv-cu128/bin/python
```

The sibling /workspace/runpod-slim/.venv-cu128 environment is also detected. This is a startup/setup bundle for an existing ComfyUI image, not a fresh OS installer. Linux needs bash, git, curl, flock, tar and python3. Expose port 8188 as HTTP.

Set **HF_TOKEN** as a private pod/template environment variable if needed, and accept the official FLUX.2 Klein 9B license with that Hugging Face account. Do not paste tokens into GitHub or this PDF.

This entrypoint starts ComfyUI only. If you need Jupyter or other image services, retain the image supervisor and have it run the GitHub bootstrap as its 8188 service.

2. Models installed automatically

MUNGBEANS / RunPod reference / Updated 22 October 2026 (Perth)

11 locked files / 112.30 GB (104.59 GiB), excluding ComfyUI, existing models and outputs. Allow at least 130 GB free for the model set, plus space for generations.

Model	Required support / loader
Phr00t Rapid v23 Q8_0 GGUF	Qwen2.5-VL 7B Q8_0 GGUF + matching BF16 mmproj; Qwen Image VAE. UnetLoaderGGUF and CLIPLoaderGGUF type qwen_image.
FLUX.2 Klein 9B distilled BF16	Qwen3 8B BF16 encoder + Flux2 VAE. Core UNETLoader; CLIPLoader type flux2. Official weights may be gated.
Miraclein NSFW v3.0 FP8 9B	Shares Klein Qwen3 8B encoder and Flux2 VAE. Core UNETLoader. Exact named file from EllaPriest45 mirror.
Aisha NSFW beta v9.7 distilled BF16	Shares Klein Qwen3 8B encoder and Flux2 VAE. Downloaded from the official Aisha dataset repository.
KreaNoct V4 INT8 ConvRot	Qwen3-VL 4B FP8-scaled encoder + shared Qwen Image VAE. Core diffusion-model loader. Publisher requires ComfyUI >=0.37.

Important workflow differences

Phr00t GGUF is loaded as a diffusion model with separate encoder and VAE. Do not use the AIO safetensors checkpoint-loading commands from the old October 13 guide for this GGUF setup. The mmproj stays beside its encoder with the exact name in the manifest.

KreaNoct: publisher recommends 8 steps, CFG 1, euler/simple. Phr00t v23: euler_ancestral/beta, CFG 1, 4-8 steps. Miraclein mirror filename suggests 12 steps/euler/CFG 1.1. Use a matching workflow for Aisha and the official distilled Klein workflow for Klein.

Custom nodes and compatibility

Installs only pinned ComfyUI-GGUF, matching the minimal recovered node installer. Existing unrelated nodes are preserved. No automatic core, torch, CUDA or existing package upgrade. Only missing small GGUF dependencies are installed without dependency resolution.

Setup checks for Klein/Qwen/Krea2 support in the existing core. Missing support is reported; the core is not silently updated. Model pairing was checked from repository records, the GGUF loader source and relevant tensor headers. Actual target-GPU inference still needs testing.

3. Check startup and recover

MUNGBEANS / RunPod reference / Updated 22 October 2026 (Perth)

Logs and serving checks

```
tail -n 100 /workspace/.runpod-startup/startup.log
cat /workspace/.runpod-startup/setup-status.json
curl --connect-timeout 5 --max-time 15 -fsS \
  http://127.0.0.1:8188/system_stats
curl --connect-timeout 5 --max-time 15 -fsS \
  http://127.0.0.1:8188/object_info/UnetLoaderGGUF
curl --connect-timeout 5 --max-time 15 -fsS \
  http://127.0.0.1:8188/object_info/CLIPLoaderGGUF
```

First boot downloads models before launching 8188; this can take a long time. A failed gated file is reported individually and the available ComfyUI installation still launches. A busy 8188 causes exit rather than killing its process. Do not launch duplicates.

Persistent files and a different GPU

Mount the same persistent storage at the expected paths to reuse models on another compatible GPU pod. A new empty disk means a new download. Network volumes must be available in the chosen GPU data center.

Completed files are verified with size and SHA256, then tracked by a fingerprint receipt. Unchanged files are not downloaded or rehashed each launch. Set **VERIFY_ALL=1** for the next normal launch to rehash everything, then remove the setting.

Interrupted transfers resume from .part files. A mismatching completed file or corrupt full partial is preserved and reported. Inspect it against models.lock.json before moving it aside manually; automatic startup does not repeatedly replace it.

Instance paths and optional settings

```
COMFY_ROOT=/workspace/runpod-slim/ComfyUI
COMFY_PYTHON=/workspace/runpod-slim/ComfyUI/.venv-cu128/bin/python
COMFY_INSTANCE_DIR=/workspace/comfy8188
STARTUP_REF=main
```

These are optional template environment settings. The instance directory contains user/input/output/temp and its own database. Point it at your existing 8188 instance directory if you want those workflows/settings. Never share it with the live 8189 instance.

GitHub bundles are cached under /workspace/.runpod-startup-source. An unavailable GitHub fetch falls back to the validated cached bundle. The first boot needs access to GitHub. STARTUP_REF can be a reviewed full Git commit SHA to pin the bundle.

Proof still required on the pod

Local checks passed Bash syntax and five real HTTP download fixtures: reuse, resume, bad-part preservation, corrupt-file rejection and HTTP 403 handling. No Linux/GPU workflow has been run for this bundle. After startup, open the 8188 proxy and generate a small test image with each model you intend to use.

4. Exact files and source record

MUNGBEANS / RunPod reference / Updated 22 October 2026 (Perth)

All files are locked to full repository commit revisions and SHA256 hashes in **models.lock.json**. The following is the download inventory; the machine-readable manifest is the authority for complete hashes.

diffusion_models/Qwen-Rapid-NSFW-v23_Q8_0.gguf

Novice25/Qwen-Image-Edit-Rapid-AIO-GGUF / 21.751 GB

text_encoders/Qwen2.5-VL-7B-Instruct-Q8_0.gguf

unsloth/Qwen2.5-VL-7B-Instruct-GGUF / 8.099 GB

text_encoders/Qwen2.5-VL-7B-Instruct-mmproj-BF16.gguf

QuantStack/Qwen-Image-Edit-GGUF / 1.354 GB

vae/qwen_image_vae.safetensors

Comfy-Org/Qwen-Image_ComfyUI / 0.254 GB

diffusion_models/flux-2-klein-9b.safetensors

black-forest-labs/FLUX.2-klein-9B / 18.157 GB

text_encoders/qwen_3_8b.safetensors

Comfy-Org/flux2-klein-9B / 16.382 GB

vae/flux2-vae.safetensors

Comfy-Org/flux2-klein-9B / 0.336 GB

diffusion_models/aisha_nsfw_beta_v9_7_distilled_bf16.safetensors

Aisha-AI-Official/flux-2-klein-models / 18.157 GB

diffusion_models/Miraclein NSFW v3.0 FP8 - Klein9B - 12steps,euler,cfg1.1.safetensors

EllaPriest45/Klein9B_Checkpoints / 9.079 GB

diffusion_models/KreaNoct_V4_int8_convrot.safetensors

Noctaluna/Krea2-Noct-Uncensored / 13.493 GB

text_encoders/qwen3vl_4b_fp8_scaled.safetensors

Comfy-Org/Krea-2 / 5.242 GB

GitHub reference

github.com/martynguyen420-debug/MUNGBEAN / runpod-startup

Files: github-bootstrap.sh, runpod-template-start.sh, setup.py, models.lock.json, nodes.lock.json, template-start.json, template-command.txt and README.md. GitHub hosts only the scripts/configuration; model bytes come from the pinned Hugging Face repositories.

Aisha source: huggingface.co/datasets/Aisha-AI-Official/flux-2-klein-models. Miraclein source: huggingface.co/EllaPriest45/Klein9B_Checkpoints, a third-party Civitai backup. KreaNoct source: huggingface.co/Noctaluna/Krea2-Noct-Uncensored. Full sources and limitations are in README.md.

Updated from the newest RunPod commands PDF found on F: (13 October 2026). Original preserved. Recovered shell scripts provided the exact Aisha/Miraclein URLs. Publisher/API metadata and relevant headers rechecked on 22 October 2026.